

# Supervised Learning Regression – Course Project

## IBM Machine Learning

Author: Isidro Brevers Gómez

E-mail: ibrevers@gmail.com

### 1. Objective and brief description of the data set and a summary of its attributes

The data set used contains information about traffic accidents in Madrid, Spain in 2018, which is published by Madrid City Council at <https://datos.madrid.es/portal/site/egob/> in CSV format.

The objective of the analysis will be developing a model that allows predicting the number of casualties resulting from accidents, based on certain parameters such as the day of the week, the hour, weather and road conditions, districts, type of accident and persons involved, and other.

The data set consists on a register for every person involved in the accident, including (non-useful data disregarded through cleansing work is omitted, full detail of the data set can be found at [https://datos.madrid.es/FWProjects/egob/Catalogo/Seguridad/Ficheros/Estructura\\_DS\\_Accidentes\\_trafico\\_2010\\_2018.pdf](https://datos.madrid.es/FWProjects/egob/Catalogo/Seguridad/Ficheros/Estructura_DS_Accidentes_trafico_2010_2018.pdf)):

- *Fecha* [object]: accident date.
- *Rango horario* [object]: hourly range when the accident took place.
- *Día semana* [object]: day of the week when the accident took place.
- *Distrito* [object]: district where the accident took place.
- *CPFA Granizo* [object]: weather conditions hail.
- *CPFA Hielo* [object]: weather conditions ice.
- *CPFA Lluvia* [object]: weather conditions rain.
- *CPFA Niebla* [object]: weather conditions fog.
- *CPFA Seco* [object]: weather conditions dry.
- *CPFA Nieve* [object]: weather conditions snow.
- *CPSV Mojada* [object]: road conditions wet.
- *CPSV Aceite* [object]: road conditions oil.
- *CPSV Barro* [object]: road conditions mud.
- *CPSV Grava Suelta* [object]: road conditions stones.
- *CPSV Hielo* [object]: road conditions ice.
- *CPSV Seca Y Limpia* [object]: road conditions dry and clean.
- *Victimas* [int64]: number of casualties in the accident.
- *TIPO ACCIDENTE* [object]: type of accident
- *Tipo Vehiculo* [object]: type of vehicle involved in the accident.
- *TIPO PERSONA* [object]: type of person involved in the accident.
- *SEXO* [object]: sex of person involved in the accident.
- *LESIVIDAD* [object]: damage degree.
- *Tramo Edad* [object]: age range of the person involved in the accident.

## 2. Initial plan for data exploration

As a first step, the CSV file was downloaded and read into a dataframe for further exploration and analysis, resulting into the following (first 5 results):

	FECHA	RANGO HORARIO	DIA SEMANA	DISTRITO	LUGAR ACCIDENTE	Nº	Nº PARTE	CPFA Granizo	CPFA Hielo	CPFA Lluvia	...	CPSV Grava Suelta	CPSV Hielo	CPSV Seca Y Limpia	* Nº VICTIMAS	TIPO ACCIDENTE	Tipo Vehículo	TIPO PERSONA	SEXO	LESIVIDAD	Tramo Edad
0	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	CALLE DE SAN BASILIO - CALLE DEL CRISTO DE LA ...	0	2018/1	NO	NO	NO	...	NO	NO	SI	1	ATROPELLO	NO ASIGNADO	PEATON	HOMBRE	HG	DE 15 A 17 AÑOS
1	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	CALLE DE SAN BASILIO - CALLE DEL CRISTO DE LA ...	0	2018/1	NO	NO	NO	...	NO	NO	SI	1	ATROPELLO	NO ASIGNADO	TESTIGO	HOMBRE	IL	DE 30 A 34 AÑOS
2	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	CALLE DE SAN BASILIO - CALLE DEL CRISTO DE LA ...	0	2018/1	NO	NO	NO	...	NO	NO	SI	1	ATROPELLO	TURISMO	CONDUCTOR	HOMBRE	IL	DE 35 A 39 AÑOS
3	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	AVENIDA DE FRANCISCO PI Y MARGALL - AVENIDA DE...	2018/3		NO	NO	NO	...	NO	NO	SI	1	CHOQUE CON OBJETO FUJO	NO ASIGNADO	TESTIGO	HOMBRE	IL	DE 21 A 24 AÑOS
4	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	AVENIDA DE FRANCISCO PI Y MARGALL - AVENIDA DE...	2018/3		NO	NO	NO	...	NO	NO	SI	1	CHOQUE CON OBJETO FUJO	NO ASIGNADO	TESTIGO	MUJER	IL	DE 40 A 44 AÑOS

The above represents al the traffic accidents occurred in Madrid, Spain in 2018, including information about whether and road conditions, when and where the accident took place, the persons and type of vehicles involved and other.

Then, some initial checks were done in order to verify:

- The size of the data: 30,122 rows, 26 columns.
- Data does not include null or inconsistent values.
- Data types.
- Basic statistical attributes of the data.

Data columns (total 23 columns):		Victimas	
FECHA	26463 non-null object	count	26463.000000
RANGO HORARIO	26463 non-null object		
DIA SEMANA	26463 non-null object		
DISTRITO	26463 non-null object		
CPFA Granizo	26463 non-null object	mean	1.488115
CPFA Hielo	26463 non-null object		
CPFA Lluvia	26463 non-null object		
CPFA Niebla	26463 non-null object	std	1.278060
CPFA Seco	26463 non-null object		
CPFA Nieve	26463 non-null object		
CPSV Mojada	26463 non-null object	min	1.000000
CPSV Aceite	26463 non-null object		
CPSV Barro	26463 non-null object		
CPSV Grava Suelta	26463 non-null object	25%	1.000000
CPSV Hielo	26463 non-null object		
CPSV Seca Y Limpia	26463 non-null object	50%	1.000000
Victimas	26463 non-null int64		
TIPO ACCIDENTE	26463 non-null object	75%	2.000000
Tipo Vehículo	26463 non-null object		
TIPO PERSONA	26463 non-null object		
SEXO	26463 non-null object		
LESIVIDAD	26463 non-null object	max	19.000000
Tramo Edad	26463 non-null object		
dtypes: int64(1), object(22)			

## 3. Actions taken for data cleansing and feature engineering

Once data was read, some data cleaning was necessary in order to ensure a better understanding and further analysis of the information.

- Some columns were dropped and renamed.
- It was check that the data set does not include null or missing values.
- Data where TIPO PERSONA was TESTIGO was dropped, as witness were not involved in the accident thus such information is not useful.


	FECHA	RANGO HORARIO	DIA SEMANA	DISTRITO	CPFA Granizo	CPFA Hielo	CPFA Lluvia	CPFA Niebla	CPFA Seco	CPFA Nieve	...	CPSV Grava Suelta	CPSV Hielo	CPSV Seca Y Limpia	Victimas	TIPO ACCIDENTE	Tipo Vehículo	TIPO PERSONA	SEXO	LESIVIDAD	Tramo Edad
0	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	NO	NO	NO	NO	SI	NO	...	NO	NO	SI	1	ATROPELLO	NO ASIGNADO	PEATON	HOMBRE	HG	DE 15 A 17 AÑOS
2	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	NO	NO	NO	NO	SI	NO	...	NO	NO	SI	1	ATROPELLO	TURISMO	CONDUCTOR	HOMBRE	IL	DE 35 A 39 AÑOS
7	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	NO	NO	NO	NO	SI	NO	...	NO	NO	SI	1	CHOQUE CON OBJETO FUJO	TURISMO	CONDUCTOR	HOMBRE	IL	DE 21 A 24 AÑOS
8	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	NO	NO	NO	NO	SI	NO	...	NO	NO	SI	1	CHOQUE CON OBJETO FUJO	TURISMO	VIAJERO	HOMBRE	HG	DE 21 A 24 AÑOS
9	01/01/2018	DE 1:00 A 1:59	LUNES	SAN BLAS	NO	NO	NO	NO	SI	NO	...	NO	NO	SI	1	COLISIÓN DOBLE	TURISMO	CONDUCTOR	HOMBRE	IL	DE 25 A 29 AÑOS

- For visualization purposes, specific dataframes were created grouping data, delete columns and setting specific indexes.

#### 4. Key Findings and Insights, which synthesizes the results of Exploratory Data Analysis in an insightful and actionable manner

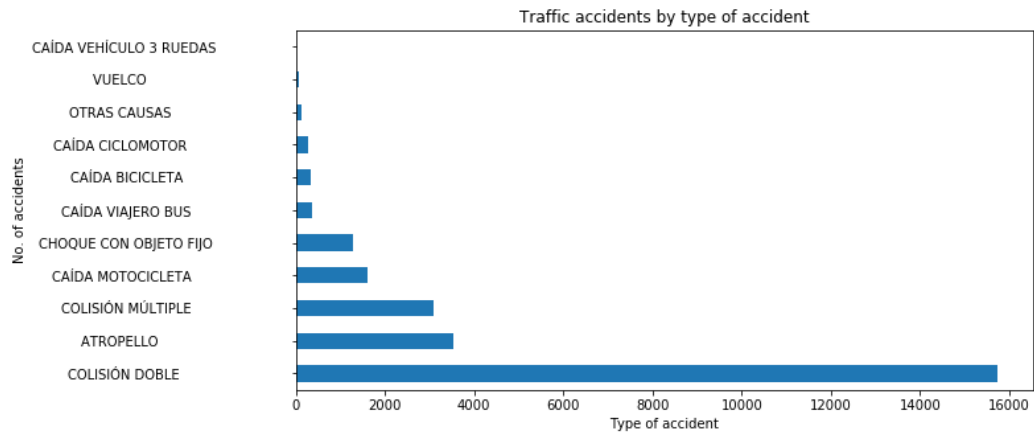
Some visualizations were plotted for the purposes of a better understanding of the information.

##### 4.1. Traffic accidents by day of the week




As per the above bar chart, we can extract that Friday is the day of the week when more accidents take place, while on the other hand weekends (Saturday and Sunday) are the days when less accidents take place.

#### 4.2. Traffic accidents by type of accident



As per the above bar chart, we can extract that most of the accidents are double collisions between vehicles, followed by pedestrian impacts and multiple collisions.

#### 4.3. Traffic accidents by type of person involved and sex




As per the above pie chart, we can extract that in most of the accidents the type of person involved was the driver, as well as that in most accidents the person involved was male.

#### 4.4. Casualties in traffic accidents by date

	FECHA	Victimas
0	07/03/2018	514
1	20/07/2018	487
2	17/07/2018	345
3	23/05/2018	286
4	15/04/2018	271
5	06/04/2018	238
6	09/05/2018	236
7	27/04/2018	226
8	06/01/2018	216
9	06/10/2018	212

## 5. Determining normality

Victimas	
count	26463.000000
mean	1.488115
std	1.278060
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	19.000000




Accidents occurred in Madrid in 2018 had a minimum number of casualties of 1 and a maximum of 19, being the mean 1.49 and the standard deviation 1.28.

Our target, the number of casualties (Victimas) does not look normal. If we run D'Agostino test, we obtain that P-value is close to 0, so our variable y is not normally distributed.


```
NormaltestResult(statistic=31902.082170149784, pvalue=0.0)
```

### Testing log / Square root / Box cox




```
normaltest(log_casualties)
```

```
NormaltestResult(statistic=10340.212264191083, pvalue=0.0)
```



```
normaltest(sqrt_casualties)
```

```
NormaltestResult(statistic=19078.37606717139, pvalue=0.0)
```



```
normaltest(boxcox_casualties)
```

```
NormaltestResult(statistic=6011.159265145325, pvalue=0.0)
```

As per the results above, it has not been possible to obtain a normally distributed y-variable.

## 6. One-hot encoding

Since all the variables except y-variable are categorical, we need first to apply one-hot encoding to the data set. For these purposes, 107 extra columns will be created, giving raise to a one-hot encoded dataframe with 129 columns:

Victimas	RANGO HORARIO_0	RANGO HORARIO_1	RANGO HORARIO_2	RANGO HORARIO_3	RANGO HORARIO_4	RANGO HORARIO_5	RANGO HORARIO_6	RANGO HORARIO_7	RANGO HORARIO_8	...	CPSV Barro_0	CPSV Barro_1	CPSV Grava Suelta_0	CPSV Grava Suelta_1	CPSV Hielo_0	CPSV Hielo_1	CPSV Seca Y Limpia_0	CPSV Seca Y Limpia_1	CPSV Aceite_0	CPSV Aceite_1
0	1	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	1.0	0.0
2	1	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	1.0	0.0
7	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	1.0	0.0
8	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	1.0	0.0
9	1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	1.0	0.0

5 rows × 129 columns

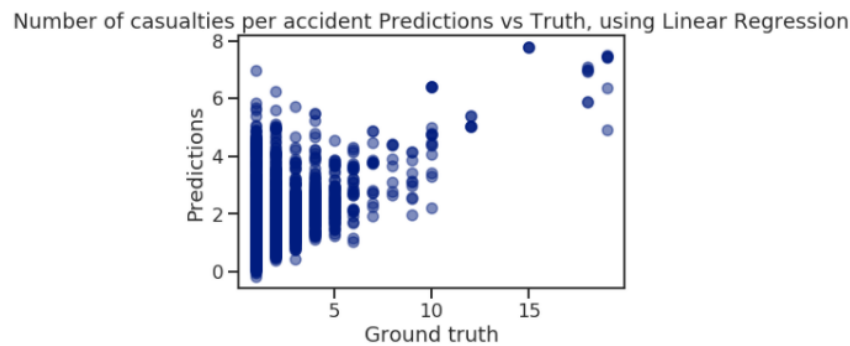
## 7. Train and test

Once the data set is ready, we will proceed to train and test, fit a basic linear regression model on the training data and calculate mean squared error on both the train and test sets.

```
[train    1.057112
 test     1.099653
 Name: one-hot enc, dtype: float64]
```

Error values are very different for the train and test data. In particular, the errors on the test data are much higher, which could be due to the one-hot encoded model being overfitting the data.

Finally, we plot predictions vs actual for the one-hot encoded model:



## 8. Cross validation

### 8.1. K-Folds

We will first split the data set into three folds using KFold object and proceed to train and test the basic linear regression model again, and calculate mean squared error:

```
[0.25680219387013503, 0.22645297494653172, 0.25643158986857195]
```

Mean squared error has significantly reduced, so our model is performing better.


If we do the same but appropriately scaling our data as we go through the folds, we obtain:

```
[0.25666920400037796, 0.22649752484204733, 0.2564449533587234]
```

Results are almost the same, because for vanilla linear regression with no regularization, scaling does not matter for performance.

## 8.2. Hyperparameter tuning

We first generate an exponentially spaced range of values using the numpy geomspace function, and then tune the alpha hyperparameter for Lasso regression:



As per the above, we have found that  $\alpha = 0.01$  would be optimum, as a higher degree would not significantly improve the model.

## 9. Suggestions for next steps in analyzing the data

In order to improve and go deeper into the analysis done, the following is proposed:

- The current dataset to be updated and complemented with additional data.
- As per the results of section 5, it has not been possible to obtain a normally distributed y-variable through Testing log / Square root / Box cox, so it is suggested to further analyze this and look for other reasons why such results are obtained and additional tools.
- Increase the number of iterations for Lasso.
- Keep training and testing the model using multiple and non-linear regression, as well as other regularization and cross validation techniques, polynomial features, ridge regression...

In [3]:

```
# Import Libraries
import pandas as pd
import numpy as np
from scipy import stats
import matplotlib.pyplot as plt
%matplotlib inline
```

In [4]:

```
# Download information (source: Madrid City Council https://datos.madrid.es)
!wget -q -O 'accidents.csv' https://datos.madrid.es/egob/catalogo/300228-18-accidentes-tra
fico-detalle.csv
```

In [5]:

```
# Check the encoding of the csv
#!pip install chardet
#import chardet

#with open("accidents.csv", 'rb') as file:
#    print(chardet.detect(file.read()))
```




In [6]:

```
df_accidents = pd.read_csv('accidents.csv', sep=';', encoding='ISO-8859-1')
df_accidents.head(5)
```

Out[6]:

	FECHA	RANGO HORARIO	DIA SEMANA	DISTRITO	LUGAR ACCIDENTE	Nº	Nº PARTE	CPFA Granizo	CPFA Hielo	CPF Lluvi
0	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	CALLE DE SAN BASILIO - CALLE DEL CRISTO DE LA ...	0	2018/1	NO	NO	Nº
1	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	CALLE DE SAN BASILIO - CALLE DEL CRISTO DE LA ...	0	2018/1	NO	NO	Nº
2	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	CALLE DE SAN BASILIO - CALLE DEL CRISTO DE LA ...	0	2018/1	NO	NO	Nº
3	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	AVENIDA DE FRANCISCO PI Y MARGALL - AVENIDA DE...		2018/3	NO	NO	Nº
4	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	AVENIDA DE FRANCISCO PI Y MARGALL - AVENIDA DE...		2018/3	NO	NO	Nº

5 rows × 26 columns



In [7]:

```
df_accidents.shape
```

Out[7]:

```
(30122, 26)
```

In [8]:

```
df_accidents.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30122 entries, 0 to 30121
Data columns (total 26 columns):
FECHA                30122 non-null object
RANGO HORARIO        30122 non-null object
DIA SEMANA           30122 non-null object
DISTRITO             30122 non-null object
LUGAR ACCIDENTE      30122 non-null object
Nº                   30122 non-null object
Nº PARTE             30122 non-null object
CPFA Granizo         30122 non-null object
CPFA Hielo           30122 non-null object
CPFA Lluvia          30122 non-null object
CPFA Niebla          30122 non-null object
CPFA Seco            30122 non-null object
CPFA Nieve           30122 non-null object
CPSV Mojada          30122 non-null object
CPSV Aceite          30122 non-null object
CPSV Barro           30122 non-null object
CPSV Grava Suelta    30122 non-null object
CPSV Hielo           30122 non-null object
CPSV Seca Y Limpia    30122 non-null object
* Nº VICTIMAS        30122 non-null int64
TIPO ACCIDENTE       30122 non-null object
Tipo Vehiculo        30122 non-null object
TIPO PERSONA         30122 non-null object
SEXO                 30122 non-null object
LESIVIDAD            30122 non-null object
Tramo Edad           30122 non-null object
dtypes: int64(1), object(25)
memory usage: 6.0+ MB
```

In [9]:

```
df_accidents.describe()
```

Out[9]:

	* Nº VICTIMAS
count	30122.000000
mean	1.459232
std	1.232216
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	19.000000

## Data cleasing

In [10]:

```
# Delete columns
df_accidents = df_accidents.drop(df_accidents.columns[[4, 5, 6]], axis=1)

# Rename columns
df_accidents.rename(columns={'* Nº VICTIMAS': 'Victimas'}, inplace=True)
```

In [11]:

```
df_accidents.shape
```

Out[11]:

```
(30122, 23)
```

In [12]:

```
# Check if there is any missing value  
df_accidents.isna().sum()
```

Out[12]:

FECHA	0
RANGO HORARIO	0
DIA SEMANA	0
DISTRITO	0
CPFA Granizo	0
CPFA Hielo	0
CPFA Lluvia	0
CPFA Niebla	0
CPFA Seco	0
CPFA Nieve	0
CPSV Mojada	0
CPSV Aceite	0
CPSV Barro	0
CPSV Grava Suelta	0
CPSV Hielo	0
CPSV Seca Y Limpia	0
Victimas	0
TIPO ACCIDENTE	0
Tipo Vehiculo	0
TIPO PERSONA	0
SEXO	0
LESIVIDAD	0
Tramo Edad	0

dtype: int64

In [13]:

```
# Check if there is any null value
df_accidents.isnull().sum()
```

Out[13]:

```
FECHA                0
RANGO HORARIO        0
DIA SEMANA           0
DISTRITO             0
CPFA Granizo         0
CPFA Hielo           0
CPFA Lluvia          0
CPFA Niebla          0
CPFA Seco            0
CPFA Nieve           0
CPSV Mojada          0
CPSV Aceite          0
CPSV Barro           0
CPSV Grava Suelta   0
CPSV Hielo           0
CPSV Seca Y Limpia   0
Victimas             0
TIPO ACCIDENTE       0
Tipo Vehiculo        0
TIPO PERSONA         0
SEXO                 0
LESIVIDAD            0
Tramo Edad           0
dtype: int64
```

In [14]:

```
# Drop rows with value TIPO PERSONA = TESTIGO, as witness were not directly involved in the accident
df_accidents = df_accidents[df_accidents['TIPO PERSONA'] != 'TESTIGO']
```

In [15]:

df\_accidents.head()

Out[15]:

	FECHA	RANGO HORARIO	DIA SEMANA	DISTRITO	CPFA Granizo	CPFA Hielo	CPFA Lluvia	CPFA Niebla	CPFA Seco	CPFA Nieve
0	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	NO	NO	NO	NO	SI	NO
2	01/01/2018	DE 00:00 A 00:59	LUNES	USERA	NO	NO	NO	NO	SI	NO
7	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	NO	NO	NO	NO	SI	NO
8	01/01/2018	DE 1:00 A 1:59	LUNES	HORTALEZA	NO	NO	NO	NO	SI	NO
9	01/01/2018	DE 1:00 A 1:59	LUNES	SAN BLAS	NO	NO	NO	NO	SI	NO

5 rows × 23 columns



## EDA

In [16]:

```
# Traffic accidents by day of the week
df_day = df_accidents['DIA SEMANA'].value_counts()
df_day
```

Out[16]:


```
VIERNES      4721
MIERCOLES    4042
MARTES       4028
LUNES        3906
JUEVES       3899
SABADO       3018
DOMINGO      2849
Name: DIA SEMANA, dtype: int64
```

In [17]:

```
df_day.plot(kind='bar',
             stacked=False,
             figsize=(10, 5))

plt.title('Traffic accidents by day of the week')
plt.ylabel('No. of accidents')
plt.xlabel('Day of the week')

plt.show()
```



In [18]:


```
# Traffic accidents by type of accident
df_tipo = df_accidents['TIPO ACCIDENTE'].value_counts()
```

In [19]:

```
df_tipo.plot(kind='barh',
             stacked=True,
             figsize=(10, 5))

plt.title('Traffic accidents by type of accident')
plt.ylabel('No. of accidents')
plt.xlabel('Type of accident')

plt.show()
```



In [20]:

```
# Traffic accidents by type of person involved
df_person = df_accidents['TIPO PERSONA'].value_counts()
df_person
```

Out[20]:


CONDUCTOR	19333
VIAJERO	5365
PEATON	1765

Name: TIPO PERSONA, dtype: int64



In [21]:

```
df_person.plot(kind='pie',  
               stacked=True,  
               figsize=(10, 5))  
  
plt.title('Traffic accidents by type of person involved')  
  
plt.show()
```



In [22]:

```
# Traffic accidents by sex  
df_sex = df_accidents['SEXO'].value_counts()  
df_sex
```

Out[22]:

HOMBRE	16556
MUJER	8682
NO ASIGNADO	1225

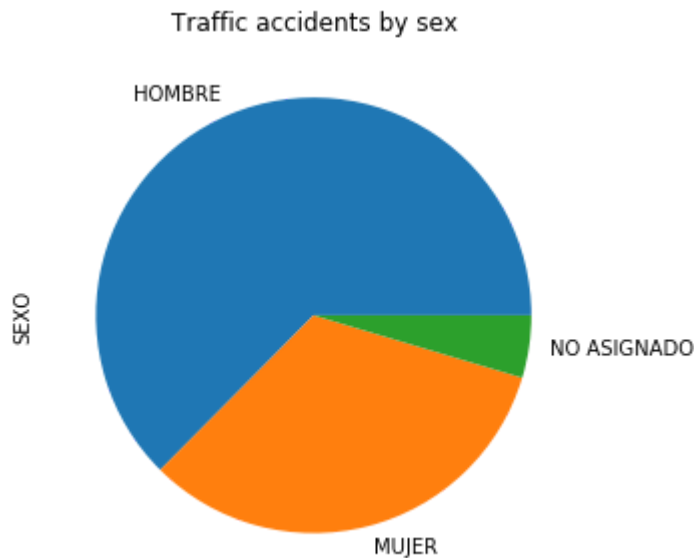
Name: SEXO, dtype: int64

In [23]:

```
df_sex.plot(kind='pie',
            stacked=True,
            figsize=(10, 5))

plt.title('Traffic accidents by sex')

plt.show()
```



In [24]:

```
# Casualties by date
df_date = df_accidents.groupby('FECHA').sum()
df_date.sort_values(['Victimas'], ascending=False, axis=0, inplace=True)
df_date.head(10).reset_index()
```

Out[24]:

	FECHA	Victimas
0	07/03/2018	514
1	20/07/2018	487
2	17/07/2018	345
3	23/05/2018	286
4	15/04/2018	271
5	06/04/2018	238
6	09/05/2018	236
7	27/04/2018	226
8	06/01/2018	216
9	06/10/2018	212

In [25]:

```
df_accidents.describe()
```


Out[25]:

Victimas	
count	26463.000000
mean	1.488115
std	1.278060
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	19.000000

## Determining normality

In [26]:

```
df_accidents.Victimas.hist();
```



In [27]:

```
from scipy.stats.mstats import normaltest # D'Agostino K^2 Test
normaltest(df_accidents.Victimas.values)
```

Out[27]:

```
NormaltestResult(statistic=31902.082170149784, pvalue=0.0)
```

In [28]:


```
# Testing log
log_casualties = np.log(df_accidents.Victimas)
```

In [29]:

```
log_casualties.hist()
```

Out[29]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f7cdd7a6350>



In [30]:

```
normaltest(log_casualties)
```

Out[30]:

NormaltestResult(statistic=10340.212264191083, pvalue=0.0)

In [31]:


```
# Box cox
from scipy.stats import boxcox
bc_result = boxcox(df_accidents.Victimas)
boxcox_casualties = bc_result[0]
lam = bc_result[1]
```

In [32]:

```
plt.hist(boxcox_casualties)
```

Out[32]:

```
(array([19470.,    0.,    0.,    0.,    0.,    0.,    0.,    0.,  
        4445., 2548.]),  
 array([0.          , 0.03015594, 0.06031189, 0.09046783, 0.12062378,  
        0.15077972, 0.18093567, 0.21109161, 0.24124756, 0.2714035 ,  
        0.30155945]),  
<a list of 10 Patch objects>)
```



In [33]:

```
normaltest(boxcox_casualties)
```

Out[33]:


```
NormaltestResult(statistic=6011.159265145325, pvalue=0.0)
```

In [34]:

```
# Square root
sqrt_casualties = np.sqrt(df_accidents.Victimas)
plt.hist(sqrt_casualties)
```

Out[34]:

```
(array([1.947e+04, 4.445e+03, 1.917e+03, 2.790e+02, 1.680e+02, 6.500e+01,
        4.600e+01, 1.500e+01, 1.700e+01, 4.100e+01]),
 array([1.          , 1.33588989, 1.67177979, 2.00766968, 2.34355958,
        2.67944947, 3.01533937, 3.35122926, 3.68711915, 4.02300905,
        4.35889894]),
 <a list of 10 Patch objects>)
```



In [35]:

```
normaltest(sqrt_casualties)
```

Out[35]:

```
NormaltestResult(statistic=19078.37606717139, pvalue=0.0)
```

## One-hot encoding

In [36]:

```
df_accidents = df_accidents.drop(df_accidents.columns[[0]], axis=1)
df_accidents.head()
```

Out[36]:

	RANGO HORARIO	DIA SEMANA	DISTRITO	CPFA Granizo	CPFA Hielo	CPFA Lluvia	CPFA Niebla	CPFA Seco	CPFA Nieve	CPSV Mojada	...
0	DE 00:00 A 00:59	LUNES	USERA	NO	NO	NO	NO	SI	NO	NO	...
2	DE 00:00 A 00:59	LUNES	USERA	NO	NO	NO	NO	SI	NO	NO	...
7	DE 1:00 A 1:59	LUNES	HORTALEZA	NO	NO	NO	NO	SI	NO	NO	...
8	DE 1:00 A 1:59	LUNES	HORTALEZA	NO	NO	NO	NO	SI	NO	NO	...
9	DE 1:00 A 1:59	LUNES	SAN BLAS	NO	NO	NO	NO	SI	NO	NO	...

5 rows × 22 columns



In [37]:

```
# Select the object (string) columns
mask = df_accidents.dtypes == np.object
categorical_cols = df_accidents.columns[mask]
```

In [38]:

```
# Determine how many extra columns would be created
num_ohc_cols = (df_accidents[categorical_cols]
                 .apply(lambda x: x.nunique())
                 .sort_values(ascending=False))

# No need to encode if there is only one value
small_num_ohc_cols = num_ohc_cols.loc[num_ohc_cols>1]

# Number of one-hot columns is one less than the number of categories
small_num_ohc_cols -= 1

# This is 107 columns, assuming the original ones are dropped.
small_num_ohc_cols.sum()
```

Out[38]:

107



In [39]:

```
# Create a new data set where all of the above categorical features will be one-hot encoded
from sklearn.preprocessing import OneHotEncoder, LabelEncoder

# Copy of the data
data_ohc = df_accidents.copy()

# The encoders
le = LabelEncoder()
ohc = OneHotEncoder()

for col in num_ohc_cols.index:

    # Integer encode the string categories
    dat = le.fit_transform(data_ohc[col]).astype(np.int)

    # Remove the original column from the dataframe
    data_ohc = data_ohc.drop(col, axis=1)

    # One hot encode the data--this returns a sparse array
    new_dat = ohc.fit_transform(dat.reshape(-1,1))

    # Create unique column names
    n_cols = new_dat.shape[1]
    col_names = ['_'.join([col, str(x)]) for x in range(n_cols)]

    # Create the new dataframe
    new_df = pd.DataFrame(new_dat.toarray(),
                          index=data_ohc.index,
                          columns=col_names)

    # Append the new data to the dataframe
    data_ohc = pd.concat([data_ohc, new_df], axis=1)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique
```

values.

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.
```

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/_e
```

ncoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/\_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/\_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/preprocessing/\_encoders.py:415: FutureWarning: The handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)], while in the future they will be determined based on the unique values.

If you want the future behaviour and silence this warning, you can specify "categories='auto'".

In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder directly.

```
warnings.warn(msg, FutureWarning)
```

In [40]:

```
# Column difference is as calculated above
data_ohc.shape[1] - df_accidents.shape[1]
```

Out[40]:

107

In [41]:

```
print(df_accidents.shape[1])

# Remove the string columns from the dataframe
df_accidents = df_accidents.drop(num_ohc_cols.index, axis=1)

print(df_accidents.shape[1])
```

22

1

In [42]:

```
# Create train and test splits of the data set
from sklearn.model_selection import train_test_split

y_col = 'Victimas'

# Split the data that is one-hot encoded
feature_cols = [x for x in data_ohc.columns if x != y_col]
X_data_ohc = data_ohc[feature_cols]
y_data_ohc = data_ohc[y_col]

X_train_ohc, X_test_ohc, y_train_ohc, y_test_ohc = train_test_split(X_data_ohc, y_data_ohc,
                                                                    test_size=0.3, random_state=42)
```


In [43]:

```
data_ohc.head(5)
```

Out[43]:

	Victimas	RANGO HORARIO_0	RANGO HORARIO_1	RANGO HORARIO_2	RANGO HORARIO_3	RANGO HORARIO_4	RANGO HORARIO_5	HO
0	1	1.0	0.0	0.0	0.0	0.0	0.0	
2	1	1.0	0.0	0.0	0.0	0.0	0.0	
7	1	0.0	0.0	0.0	0.0	0.0	0.0	
8	1	0.0	0.0	0.0	0.0	0.0	0.0	
9	1	0.0	0.0	0.0	0.0	0.0	0.0	

5 rows × 129 columns



In [44]:

```
data_ohc.shape
```

Out[44]:

(26463, 129)

In [45]:

```
data_ohc.dtypes
```

Out[45]:

Victimas	int64
RANGO HORARIO_0	float64
RANGO HORARIO_1	float64
RANGO HORARIO_2	float64
RANGO HORARIO_3	float64
RANGO HORARIO_4	float64
RANGO HORARIO_5	float64
RANGO HORARIO_6	float64
RANGO HORARIO_7	float64
RANGO HORARIO_8	float64
RANGO HORARIO_9	float64
RANGO HORARIO_10	float64
RANGO HORARIO_11	float64
RANGO HORARIO_12	float64
RANGO HORARIO_13	float64
RANGO HORARIO_14	float64
RANGO HORARIO_15	float64
RANGO HORARIO_16	float64
RANGO HORARIO_17	float64
RANGO HORARIO_18	float64
RANGO HORARIO_19	float64
RANGO HORARIO_20	float64
RANGO HORARIO_21	float64
RANGO HORARIO_22	float64
RANGO HORARIO_23	float64
DISTRITO_0	float64
DISTRITO_1	float64
DISTRITO_2	float64
DISTRITO_3	float64
DISTRITO_4	float64
...	
TIPO PERSONA_0	float64
TIPO PERSONA_1	float64
TIPO PERSONA_2	float64
SEXO_0	float64
SEXO_1	float64
SEXO_2	float64
CPFA Niebla_0	float64
CPFA Niebla_1	float64
CPFA Granizo_0	float64
CPFA Granizo_1	float64
CPFA Hielo_0	float64
CPFA Hielo_1	float64
CPFA Lluvia_0	float64
CPFA Lluvia_1	float64
CPFA Nieve_0	float64
CPFA Nieve_1	float64
CPFA Seco_0	float64
CPFA Seco_1	float64
CPSV Mojada_0	float64
CPSV Mojada_1	float64
CPSV Barro_0	float64
CPSV Barro_1	float64
CPSV Grava Suelta_0	float64
CPSV Grava Suelta_1	float64



```
CPSV Hielo_0          float64
CPSV Hielo_1          float64
CPSV Seca Y Limpia_0  float64
CPSV Seca Y Limpia_1  float64
CPSV Aceite_0         float64
CPSV Aceite_1         float64
Length: 129, dtype: object
```

## Train and test

In [46]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

LR = LinearRegression()

# Storage for error values
error_df = list()

LR = LR.fit(X_train_ohc, y_train_ohc)
y_train_ohc_pred = LR.predict(X_train_ohc)
y_test_ohc_pred = LR.predict(X_test_ohc)

error_df.append(pd.Series({'train': mean_squared_error(y_train_ohc, y_train_ohc_pred),
                        'test' : mean_squared_error(y_test_ohc, y_test_ohc_pred)},
                        name='one-hot enc'))

error_df
```

Out[46]:

```
[train    1.221186
 test     1.194478
 Name: one-hot enc, dtype: float64]
```

In [47]:

```

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline


sns.set_context('talk')
sns.set_style('ticks')
sns.set_palette('dark')

ax = plt.axes()
# we are going to use y_test, y_test_pred
ax.scatter(y_test_ohc, y_test_ohc_pred, alpha=.5)

ax.set(xlabel='Ground truth',
       ylabel='Predictions',
       title='Number of casualties per accident Predictions vs Truth, using Linear Regression');

```

Number of casualties per accident Predictions vs Truth, using Linear Regression



## Cross validation

In [48]:

```

from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.model_selection import KFold, cross_val_predict
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.metrics import r2_score
from sklearn.pipeline import Pipeline

```

In [49]:

```

data_ohc = data_ohc.reset_index()
X = data_ohc.drop('Victimas', axis=1)
y = data_ohc.Victimas

```

In [50]:

```
kf = KFold(shuffle=True, random_state=72018, n_splits=3)
```

In [51]:

```
for train_index, test_index in kf.split(X):  
    print("Train index:", train_index[:10], len(train_index))  
    print("Test index:", test_index[:10], len(test_index))  
    print('')
```

Train index: [ 0 2 4 5 6 8 10 11 12 13] 17642

Test index: [ 1 3 7 9 14 17 18 19 20 26] 8821

Train index: [ 0 1 2 3 4 7 8 9 10 11] 17642

Test index: [ 5 6 16 22 23 24 29 38 41 43] 8821

Train index: [ 1 3 5 6 7 9 14 16 17 18] 17642

Test index: [ 0 2 4 8 10 11 12 13 15 21] 8821

In [52]:

```
scores = []  
lr = LinearRegression()  
  
for train_index, test_index in kf.split(X):  
    X_train, X_test, y_train, y_test = (X.iloc[train_index, :],  
                                       X.iloc[test_index, :],  
                                       y[train_index],  
                                       y[test_index])  
  
    lr.fit(X_train, y_train)  
  
    y_pred = lr.predict(X_test)  
  
    score = r2_score(y_test.values, y_pred)  
  
    scores.append(score)  
  
scores
```

Out[52]:

```
[0.25680219387013503, 0.22645297494653172, 0.25643158986857195]
```

In [53]:

```
# Scaling
scores = []

lr = LinearRegression()
s = StandardScaler()

for train_index, test_index in kf.split(X):
    X_train, X_test, y_train, y_test = (X.iloc[train_index, :],
                                       X.iloc[test_index, :],
                                       y[train_index],
                                       y[test_index])

    X_train_s = s.fit_transform(X_train)

    lr.fit(X_train_s, y_train)

    X_test_s = s.transform(X_test)

    y_pred = lr.predict(X_test_s)

    score = r2_score(y_test.values, y_pred)

    scores.append(score)

scores
```

Out[53]:

```
[0.25666920400037796, 0.22649752484204733, 0.2564449533587234]
```

In [54]:

```
# Hyperparameter tuning
alphas = np.geomspace(1e-9, 1e0, num=10)
alphas
```

Out[54]:

```
array([1.e-09, 1.e-08, 1.e-07, 1.e-06, 1.e-05, 1.e-04, 1.e-03, 1.e-02,
       1.e-01, 1.e+00])
```

In [55]:

```
scores = []
coefs = []
for alpha in alphas:
    las = Lasso(alpha=alpha, max_iter=100)

    estimator = Pipeline([
        ("scaler", s),
        ("lasso_regression", las)])

    predictions = cross_val_predict(estimator, X, y, cv = kf)

    score = r2_score(y, predictions)

    scores.append(score)
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 10446.1010953826
55, tolerance: 2.8133928352794486
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 11724.1566686666
03, tolerance: 3.1978390261875114
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 9784.53904647412
7, tolerance: 2.6332655991384213
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 10445.8183195387
85, tolerance: 2.8133928352794486
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 11723.7337941942
24, tolerance: 3.1978390261875114
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 9784.0540256503
9, tolerance: 2.6332655991384213
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 10442.9907485589
57, tolerance: 2.8133928352794486
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 11719.5053978885
12, tolerance: 3.1978390261875114
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 9779.2046326504
8, tolerance: 2.6332655991384213
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 10414.7337085554
85, tolerance: 2.8133928352794486
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 11677.2561621075
24, tolerance: 3.1978390261875114
    positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
```

```
ight want to increase the number of iterations. Duality gap: 9730.79217907072
6, tolerance: 2.6332655991384213
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 10134.0450643680
2, tolerance: 2.8133928352794486
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 11258.2439434733
86, tolerance: 3.1978390261875114
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 9254.69691068688
7, tolerance: 2.6332655991384213
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 7511.08337224725
6, tolerance: 2.8133928352794486
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 7417.12774263600
4, tolerance: 3.1978390261875114
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 5249.16307517615
6, tolerance: 2.6332655991384213
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 196.613746719460
32, tolerance: 2.8133928352794486
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 216.713160530391
18, tolerance: 3.1978390261875114
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 54.6670984690281
3, tolerance: 2.6332655991384213
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 2.81987916535217
66, tolerance: 2.8133928352794486
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 15.9600125718534
4, tolerance: 3.1978390261875114
```

```
positive)
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 5.8024500819774
4, tolerance: 2.6332655991384213
positive)
```

In [56]:

```
list(zip(alphas,scores))
```

Out[56]:

```
[(1e-09, 0.24887118802440888),
 (1e-08, 0.24887118886983872),
 (1e-07, 0.24887119732212748),
 (1e-06, 0.2488712816618286),
 (1e-05, 0.2488723038377859),
 (0.0001, 0.24887619113517923),
 (0.001, 0.24887174362252773),
 (0.01, 0.24633063661194332),
 (0.1, 0.1788869330665105),
 (1.0, -0.00017640843345367863)]
```



In [57]:

```
Lasso(alpha=1e-6).fit(X, y).coef_
```

```
/srv/conda/envs/notebook/lib/python3.7/site-packages/sklearn/linear_model/coo
rdinate_descent.py:475: ConvergenceWarning: Objective did not converge. You m
ight want to increase the number of iterations. Duality gap: 5449.19880610749
35, tolerance: 4.322401231908705
positive)
```

Out[57]:

```
array([ 1.36341033e-06,  1.05666359e-01,  1.89363047e-03, -1.19331324e-01,
        -3.90787029e-02, -1.20170343e-01,  9.33688623e-02,  4.56393768e-02,
        -9.64661739e-02,  6.05335791e-02, -5.52274578e-02, -4.87164983e-02,
         8.61944334e-02,  1.23820276e-02,  1.58445942e-01,  4.76525392e-03,
        -4.15339221e-02,  6.78115921e-01, -1.01043396e-01,  4.58220940e-03,
         1.56892609e-01,  2.25412728e-01, -1.11824921e-01, -7.72288935e-02,
        -1.58455522e-01,  1.09692472e-02, -7.08815978e-02,  1.01549916e-01,
         3.99628916e-02, -7.75539642e-02, -3.40424699e-03, -6.36147745e-02,
        -7.96436127e-02, -8.19901828e-02, -1.25652587e-02,  2.16759315e-02,
        -2.87228629e-02,  1.82118481e-01, -1.34283539e-01,  6.09983976e-02,
         1.42108063e-02,  6.71501452e-02,  1.36944544e-02, -4.84995060e-02,
        -3.71243868e-02, -1.98084833e-02,  1.23822641e-01,  1.07644366e-01,
         1.45941880e-01,  1.22605201e-01,  3.60671312e-02,  1.37842624e-02,
         4.90262060e-04, -5.37269708e-03, -9.64365647e-03,  1.86330345e-03,
         1.66149595e-02, -9.92548779e-04,  4.53747199e-01,  5.22706956e-03,
        -2.79796855e-02,  5.77146711e-03, -1.52774779e-01, -1.98721321e-01,
         3.70356813e-01,  1.25506173e-01,  1.65767781e+00, -4.53421490e-01,
         1.87770595e-01, -6.27939910e-01,  7.23441863e-02, -5.71800896e-01,
        -2.72989511e-01,  8.96586476e-02,  5.93589809e-02, -3.78795270e-01,
        -2.05679760e-01, -3.81182377e-01, -2.09134518e-01, -1.97566948e-01,
        -1.12014425e-01,  7.74212182e-01, -3.38581687e-01, -8.34724717e-02,
         2.77935955e-01, -4.39481832e-01, -5.27817015e-01,  1.13705155e-01,
        -1.03237781e-01, -4.71440561e-02, -2.33480722e-02,  0.00000000e+00,
         3.66315582e-02,  1.81012968e-02,  6.91560234e-01,  6.53860836e-01,
        -9.68312032e-02,  9.45224735e-01,  1.81559981e-01, -3.28899727e-01,
        -4.60143661e-01,  1.44902762e-01, -1.24731028e-02, -3.26815897e-02,
         1.74596314e-01,  2.77190973e-01, -8.92972255e-15,  4.07421817e-02,
        -1.42622396e-14,  5.94159282e-01, -5.66233171e-14,  5.87236634e-01,
        -0.00000000e+00,  3.96909696e-01, -2.93573144e-15,  3.97649520e-01,
        -0.00000000e+00, -1.11109756e-01,  0.00000000e+00, -2.02544841e-01,
         1.00690060e-16, -1.88302099e-01,  9.65636074e-16, -1.82787680e-01,
         0.00000000e+00, -5.65930170e-03,  8.75726217e-19, -1.36780829e-01,
         1.65170311e-14])
```

In [58]:


```
Lasso(alpha=1.0).fit(X, y).coef_
```

Out[58]:

```
array([ 8.05734695e-07,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00, -0.00000000e+00,  0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
         0.00000000e+00, -0.00000000e+00,  0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
         0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00, -0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00, -0.00000000e+00,  0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
         0.00000000e+00,  0.00000000e+00,  0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00,  0.00000000e+00,  0.00000000e+00,
         0.00000000e+00,  0.00000000e+00,  0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00, -0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00, -0.00000000e+00,  0.00000000e+00, -0.00000000e+00,
         0.00000000e+00,  0.00000000e+00,  0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00, -0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
         0.00000000e+00,  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
         0.00000000e+00, -0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00, -0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
         0.00000000e+00,  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
         0.00000000e+00, -0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
         0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
         0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
         0.00000000e+00, -0.00000000e+00,  0.00000000e+00,  0.00000000e+00,
        -0.00000000e+00])
```

In [59]:

```
plt.figure(figsize=(10,6))
plt.semilogx(alphas, scores, '-o')
plt.xlabel('$\\alpha$')
plt.ylabel('$R^2$');
```



In [ ]:

```
# PolynomialFeatures

pf = PolynomialFeatures(degree=3)

scores = []
alphas = np.geomspace(0.06, 6.0, 20)
for alpha in alphas:
    las = Lasso(alpha=alpha, max_iter=100)

    estimator = Pipeline([
        ("scaler", s),
        ("make_higher_degree", pf),
        ("lasso_regression", las)])

    predictions = cross_val_predict(estimator, X, y, cv = kf)

    score = r2_score(y, predictions)

    scores.append(score)
```

In [ ]:

```
plt.semilogx(alphas, scores);
```

In [ ]:

```
# Once we have found the hyperparameter (alpha~1e-2=0.01)
# make the model and train it on ALL the data
# Then release it into the wild .....
best_estimator = Pipeline([
    ("scaler", s),
    ("make_higher_degree", PolynomialFeatures(degree=2)),
    ("lasso_regression", Lasso(alpha=0.03))])

best_estimator.fit(X, y)
best_estimator.score(X, y)
```

In [ ]: